


# DATABASE TECH TALKS

## Vaccination Database Tech Talks

- Mondays @ 4:30pm (starting today)
- <https://db.cs.cmu.edu/seminar2022>



# DISK-ORIENTED ARCHITECTURE

---

The DBMS assumes that the primary storage location of the database is on non-volatile disk.

The DBMS's components manage the movement of data between non-volatile and volatile storage.

# PAGE-ORIENTED ARCHITECTURE

---

## **Insert a new tuple:**

- Check page directory to find a page with a free slot.
- Retrieve the page from disk (if not in memory).
- Check slot array to find empty space in page that will fit.

## **Update an existing tuple using its record id:**

- Check page directory to find location of page.
- Retrieve the page from disk (if not in memory).
- Find offset in page using slot array.
- Overwrite existing data (if new data fits).

# DISCUSSION

---

**What are some potential problems with the slotted page design?**

- Fragmentation
- Useless Disk I/O
- Random Disk I/O (e.g., update 20 tuples on 20 pages)

**What if the DBMS could not overwrite data in pages and could only create new pages?**

- Examples: Cloud storage (S3), HDFS

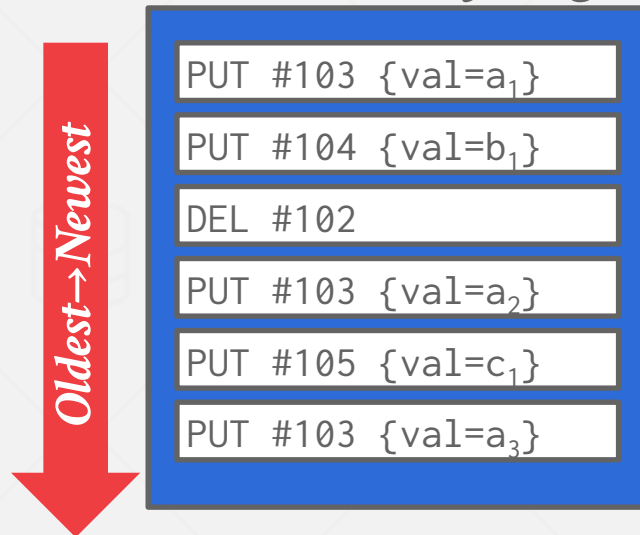
# LOG-STRUCTURED STORAGE

DBMS stores log records that contain changes to tuples (**PUT**, **DELETE**).

- Each log record must contain the tuple's unique identifier.
- Put records contain the tuple contents.
- Deletes marks the tuple as deleted.

As the application makes changes to the database, the DBMS appends log records to the end of the file without checking previous log records.

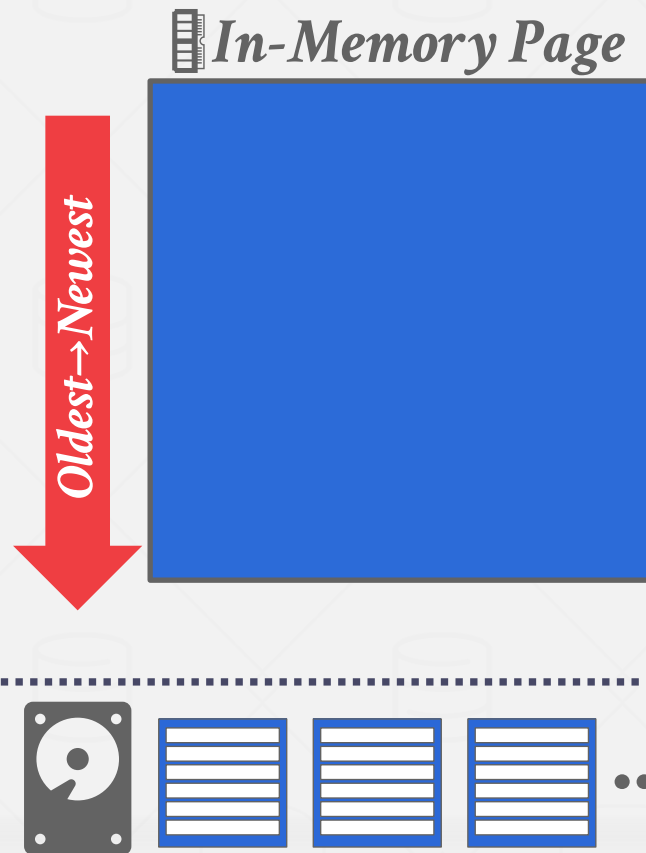
## In-Memory Page



# LOG-STRUCTURED STORAGE

When the page gets full, the DBMS writes it out disk and starts filling up the next page with records.

- All disk writes are sequential.
- On-disk pages are immutable.



# LOG-STRUCTURED STORAGE

To read a tuple with a given id, the DBMS finds the newest log record corresponding to that id.

→ Scan log from newest to oldest.


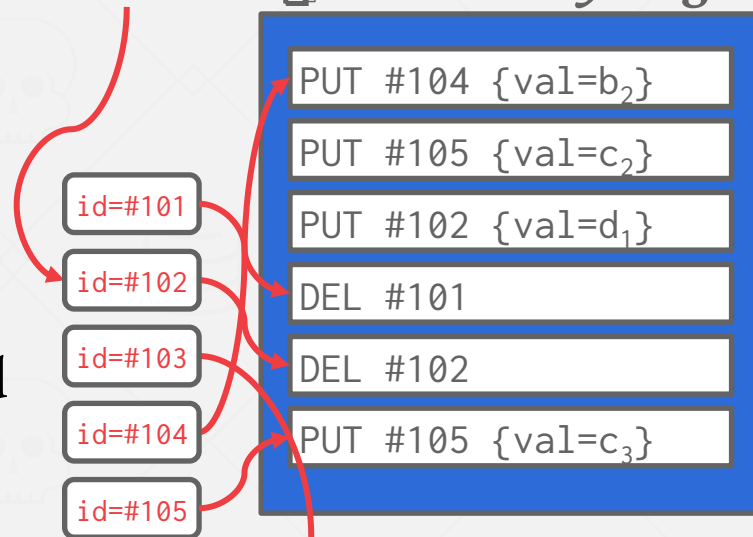
Maintain an index that maps a tuple id to the newest log record.

→ If log record is in-memory, just read it.

→ If log record is on a disk page, retrieve it.

→ We will discuss indexes in two weeks.

Get Id #102  *In-Memory Page*



# LOG-STRUCTURED COMPACTION

The log will grow forever. The DBMS needs to periodically compact pages to reduce wasted space.

*Page 1*

PUT #103 {val=a <sub>1</sub> }
PUT #104 {val=b <sub>1</sub> }
DEL #102
PUT #103 {val=a <sub>2</sub> }
PUT #105 {val=c <sub>1</sub> }
PUT #103 {val=a <sub>3</sub> }



*Page 2*

PUT #104 {val=b <sub>2</sub> }
PUT #105 {val=c <sub>2</sub> }
PUT #102 {val=d <sub>1</sub> }
DEL #101
DEL #102
PUT #105 {val=c <sub>3</sub> }



PUT #103 {val=a <sub>3</sub> }
PUT #104 {val=b <sub>2</sub> }
DEL #101
DEL #102
PUT #105 {val=c <sub>3</sub> }



# LOG-STRUCTURED COMPACTION

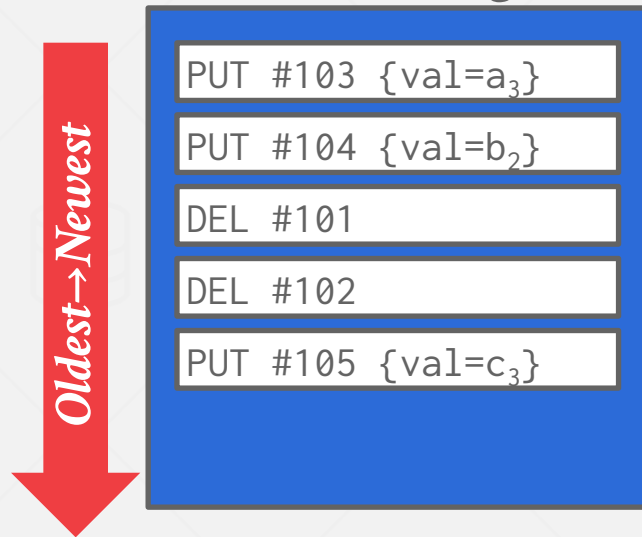
After a page is compacted, the DBMS does need to maintain temporal ordering of records within the page.

→ Each tuple id is guaranteed to appear at most once in the page.

The DBMS can instead sort the page based on id order to improve efficiency of future look-ups.

→ Called Sorted String Tables (SSTables)

## *Disk Page*




# LOG-STRUCTURED COMPACTION

After a page is compacted, the DBMS does need to maintain temporal ordering of records within the page.

→ Each tuple id is guaranteed to appear at most once in the page.

The DBMS can instead sort the page based on id order to improve efficiency of future look-ups.


→ Called Sorted String Tables (SSTables)



# LOG-STRUCTURED COMPACTION

Compaction coalesces larger log files into smaller files by removing unnecessary records.


## *Universal Compaction*



# LOG-STRUCTURED COMPACTION

Compaction coalesces larger log files into smaller files by removing unnecessary records.

## *Universal Compaction*




## *Level Compaction*

Level 0

Level 1

Level 2



# DISCUSSION

---

Log-structured storage managers are more common today. This is partly due to the proliferation of RocksDB.

**What are some downsides of this approach?**

- Write-Amplification
- Compaction is Expensive

