

Revisiting Data Prefetching for Database Systems with Machine Learning Techniques

Yu Chen, Yong Zhang, Jiacheng Wu, Jin Wang, Chunxiao Xing

Presenter: Xiling Li

Northwestern | McCORMICK SCHOOL OF
ENGINEERING

Overview

- Access page in memory is faster than access in the disk.
- Prefetch (Predict + fetch):
 - Predict future page access patterns
 - Fetch pages into buffer pool ahead before being accessed.
 - Existing methods are mostly heuristic-based methods
 - E.g. One Block Lookahead (OBL)
 - Inefficiently prefetch pages with random access by indexing
- Machine Learning (neural-network based approach)
 - Learning page access patterns from history
 - Formalize prefetching as a classification problem

Formulation

- Input:
 - A sequence of pages previously accessed
 - Represent as a sequence of page offsets (d_1, d_2, \dots, d_n)
- Output:
 - A page in the disk
 - Page offset d_t
- Classification
 - Get output distribution by softmax function

Sparsity Issue

- Sparsity of target space
 - E.g. TPC-H approximately contains 10^7 target pages
 - **Bad: Locate a page by its universal id (di)**
- Extent in MySQL can be a feasible solution
 - An extent is a group of consecutive pages within a data file
 - # of extents is much smaller than # of pages
 - **Better: Locate a page by extent it belongs and the offset in the extent**

Architecture

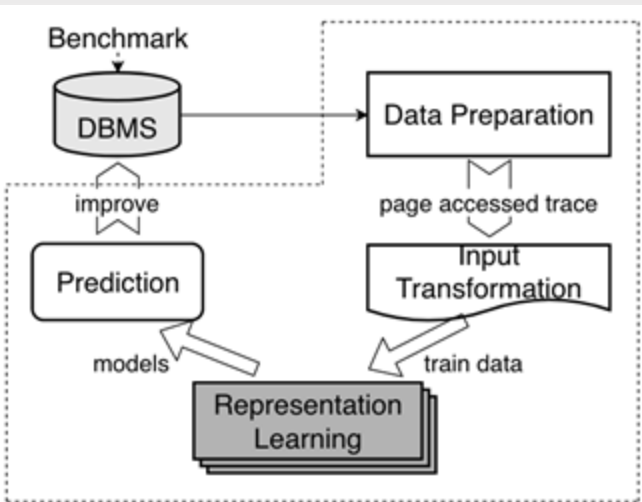


Fig. 1: Overall Architecture

- Data preparation
 - Collect access traces by running some queries
- Input transformation
 - Transform page offsets d_i into extent
 - **Overcome sparsity issue**
- Representation learning
 - Learn access pattern by traces
- Prediction
 - Fetch corresponding pages from disk

Transformation

- Transform a page offset d_i into extent offset e_i and in-extent offset f_i

$$e_i = \lfloor \frac{d_i}{\# \text{ pages in a extent}} \rfloor \quad (1)$$

$$f_i = d_i - e_i \times \# \text{ pages in a extent} \quad (2)$$

- In MySQL, each extent has 64 pages
 - E.g. if $d_i = 76$, $e_i = 1$ and $f_i = 12$
- Triplet $t_i = \langle d_i, e_i, f_i \rangle$ becomes input of latter NN models.

Single-model Framework

- Prediction model
 - Convolutional Neural Network
 - Recurrent Neural Network
 - Long Short-Term Memory
- Input: a sequence of t_i
- Output: $o_n = \{e_n, f_n\}$
- Loss function
 - Sum of two cross-entropies

$$\mathcal{L}(p(o_n), p(\hat{o}_n)) = - \sum_{e_n} p(\hat{e}_n) \log p(e_n) - \sum_{f_n} p(\hat{f}_n) \log p(f_n)$$

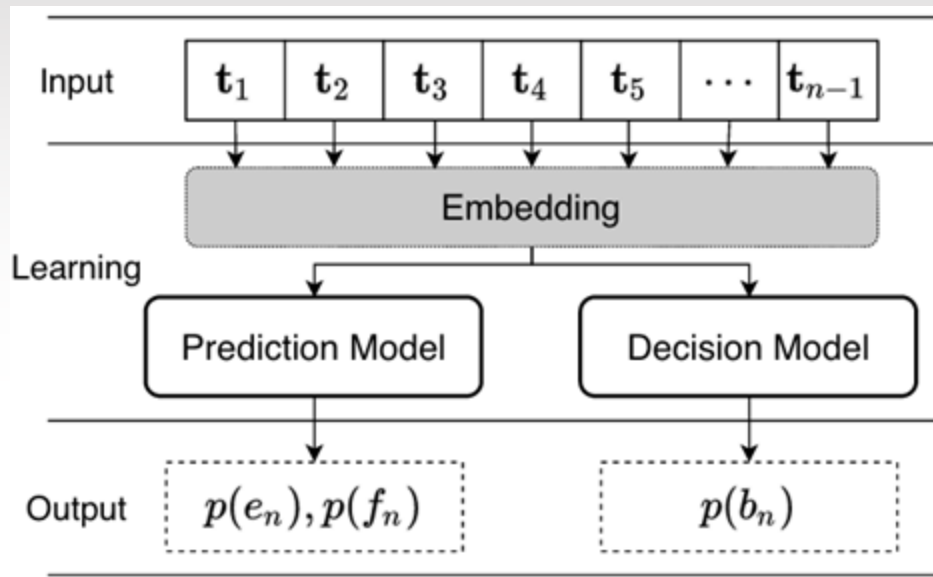


Fig. 2: Single Model with Prefetching Decision

Single-model Framework - Cont'd

- Avoiding wrong prefetching is to prevent extra disk I/O.
- Use decision model to decide if DBMS needs fetch or not at a certain timestamp n
- Input: a sequence of t_i
- Output: $b_n = \{0,1\}$
- Predict K pages for decision

Algorithm 1: Training for Decision Model

Input: Prediction Model \mathcal{M} , Input Representation t_i .

Output: Decision Model \mathcal{D} .

```
1  $train\_set = \emptyset$ .
2 foreach timestep  $n$  do
3   Obtain the label  $\hat{o}_n$ .
4   Obtain  $p(o_n)$  by applying  $\mathcal{M}$  on input series.
5   Select  $K$  page offsets from  $p(o_n)$  as collection  $\mathcal{S}$ .
6   if  $\hat{o}_n \in \mathcal{S}$  then
7     | Mark timestep  $n$  with positive label.
8   else
9     | Mark timestep  $n$  with negative label.
10  | Add timestep  $n$  into  $train\_set$ .
11 Train Decision Model  $\mathcal{D}$  on  $train\_set$ .
12 return  $\mathcal{D}$ .
```

Multi-model Framework

- TXNs from different SQL templates produce distinct access patterns
- Same architecture with different parameters for prediction models
- Classification model provides probability for each model to be chosen
- Output distribution $\{\bar{p}(e_n), \bar{p}(f_n)\} = \{\sum_{j=1}^G c_j \cdot p_j(e_n), \sum_{j=1}^G c_j \cdot p_j(f_n)\}$
- Same loss function

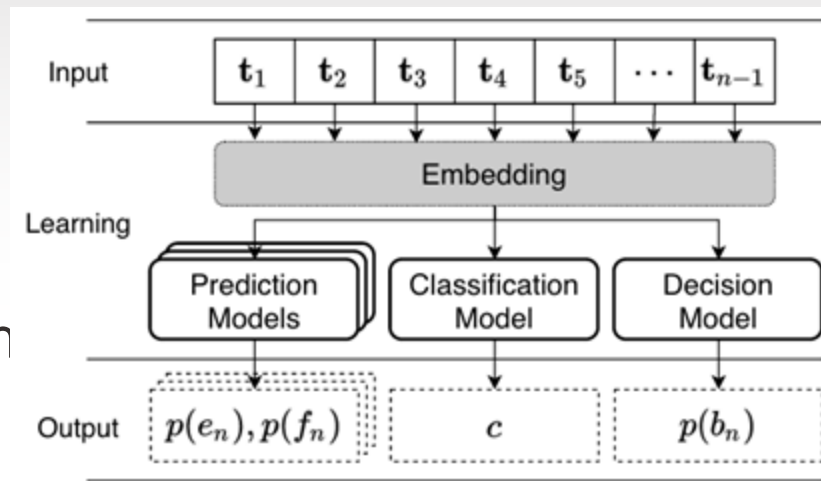


Fig. 3: Multi-Model Framework

Evaluation

TABLE I: Model Comparison

	TPC-H		TPC-DS		SSB	
	Precision(%) ¹	Recall(%)	Precision(%)	Recall(%)	Precision(%)	Recall(%)
LookAhead	20/-	81	7/-	81	22/-	88
Random	14/28	80	7/15	78	40/55	85
DNN	41/75	77	33/78	79	62/69	82
CNN	41/85	70	40/79	81	80/74	86
RNN	33/64	63	29/62	70	46/62	73
LSTM	33/64	63	30/62	71	46/63	73
Multi-Model	76/87	82	78/87	94	87/84	94

¹ We measure 2 precisions: Overall-Precision / Decision-Precision. Overall-Precision means we measure the prefetcher's precision all the timesteps, and Decision-Precision only counts when the prefetcher decides to make a prefetching.

Thank you!